

# Inference in Bayesian Networks

Andrea Passerini  
passerini@disi.unitn.it

Machine Learning

# Inference in graphical models

## Description

- Assume we have evidence **e** on the state of a subset of variables **E** in the model (i.e. Bayesian Network)
- Inference amounts at computing the posterior probability of a subset **X** of the non-observed variables given the observations:

$$p(\mathbf{X}|\mathbf{E} = \mathbf{e})$$

## Note

- When we need to distinguish between variables and their values, we will indicate random variables with uppercase letters, and their values with lowercase ones.

## Efficiency

- We can always compute the posterior probability as the ratio of two joint probabilities:

$$p(\mathbf{X}|\mathbf{E} = \mathbf{e}) = \frac{p(\mathbf{X}, \mathbf{E} = \mathbf{e})}{p(\mathbf{E} = \mathbf{e})}$$

- The problem consists of estimating such joint probabilities when dealing with a large number of variables
- Directly working on the full joint probabilities requires time exponential in the number of variables
- For instance, if all  $N$  variables are discrete and take one of  $K$  possible values, a joint probability table has  $K^N$  entries
- We would like to exploit the structure in graphical models to do inference more efficiently.

# Inference in graphical models



## Inference on a chain (1)

$$p(\mathbf{X}) = p(X_1)p(X_2|X_1)p(X_3|X_2)\cdots p(X_N|X_{N-1})$$

- The marginal probability of an arbitrary  $X_n$  is:

$$p(X_n) = \sum_{X_1} \sum_{X_2} \cdots \sum_{X_{n-1}} \sum_{X_{n+1}} \cdots \sum_{X_N} p(\mathbf{X})$$

- Only the  $p(X_N|X_{N-1})$  is involved in the last summation which can be computed first, giving a function of  $X_{N-1}$ :

$$\mu_\beta(X_{N-1}) = \sum_{X_N} p(X_N|X_{N-1})$$

# Inference in graphical models



## Inference on a chain (2)

- the marginalization can be iterated as:

$$\mu_{\beta}(X_{N-2}) = \sum_{X_{N-1}} p(X_{N-1}|X_{N-2})\mu_{\beta}(X_{N-1})$$

down to the desired variable  $X_n$ , giving:

$$\mu_{\beta}(X_n) = \sum_{X_{n+1}} p(X_{n+1}|X_n)\mu_{\beta}(X_{n+1})$$

# Inference in graphical models



## Inference on a chain (3)

- The same procedure can be applied starting from the other end of the chain, giving:


$$\mu_{\alpha}(X_2) = \sum_{X_1} p(X_1)p(X_2|X_1)$$

up to  $\mu_{\alpha}(X_n)$

- The marginal probability is now computed as the product of the contributions coming from both ends:

$$p(X_n) = \mu_{\alpha}(X_n)\mu_{\beta}(X_n)$$

# Inference in graphical models



## Inference as message passing

- We can think of  $\mu_\alpha(X_n)$  as a message passing from  $X_{n-1}$  to  $X_n$

$$\mu_\alpha(X_n) = \sum_{X_{n-1}} p(X_n | X_{n-1}) \mu_\alpha(X_{n-1})$$

- We can think of  $\mu_\beta(X_n)$  as a message passing from  $X_{n+1}$  to  $X_n$

$$\mu_\beta(X_n) = \sum_{X_{n+1}} p(X_{n+1} | X_n) \mu_\beta(X_{n+1})$$

- Each outgoing message is obtained multiplying the incoming message by the “local” probability, and summing over the node values

# Inference in graphical models

## Full message passing

- Suppose we want to know marginal probabilities for a number of different variables  $X_i$ :
  - 1 We send a message from  $\mu_\alpha(X_1)$  up to  $\mu_\alpha(X_N)$
  - 2 We send a message from  $\mu_\beta(X_N)$  down to  $\mu_\beta(X_1)$
- If all nodes store messages, we can compute any marginal probability as

$$p(X_i) = \mu_\alpha(X_i)\mu_\beta(X_i)$$

for any  $i$  having sent just a double number of messages wrt a single marginal computation



# Inference in graphical models

## Adding evidence

- If some nodes  $\mathbf{X}_e$  are observed, we simply use their observed values instead of summing over all possible values when computing their messages

## Example

$$p(\mathbf{X}) = p(X_1)p(X_2|X_1)p(X_3|X_2)p(X_4|X_3)$$

- The marginal probability of  $X_2$  and observations  $X_1 = x_{e_1}$  and  $X_3 = x_{e_3}$  is:


$$p(X_2, X_1 = x_{e_1}, X_3 = x_{e_3}) = p(X_1 = x_{e_1})p(X_2|X_1 = x_{e_1}) \cdot p(X_3 = x_{e_3}|X_2) \sum_{X_4} p(X_4|X_3 = x_{e_3})$$

## Computing conditional probability given evidence


- When adding evidence, the message passing procedure computes the joint probability of the variable *and* the evidence, and it has to be normalized to obtain the conditional probability *given* the evidence:

$$p(X_n | \mathbf{X}_e = \mathbf{x}_e) = \frac{p(X_n, \mathbf{X}_e = \mathbf{x}_e)}{\sum_{X_n} p(X_n, \mathbf{X}_e = \mathbf{x}_e)}$$


# Inference



(a)



(b)



(c)

## Inference on trees

- Efficient inference can be computed for the broadened family of tree-structured models:

**undirected trees (a)** undirected graphs with a single path for each pair of nodes


**directed trees (b)** directed graphs with a single node (the root) with no parents, and all other nodes with a single parent

**directed polytrees (c)** directed graphs with multiple parents for node and multiple roots, but still a single (undirected) path between each pair of nodes


## Description

- Efficient inference algorithms can be better explained using an alternative graphical representation called *factor graph*
- A *factor graph* is a graphical representation of a graphical model highlighting its factorization (i.e. conditional probabilities)
- The factor graph has one node for each node in the original graph
- The factor graph has one additional node (of a different type) for each factor
- A factor node has undirected links to each of the node variables in the factor


# Factor graphs: examples



$$p(x_3|x_1, x_2)p(x_1)p(x_2)$$



$$f(x_1, x_2, x_3) = p(x_3|x_1, x_2)p(x_1)p(x_2)$$



$$f_c(x_1, x_2, x_3) = p(x_3|x_1, x_2)$$

$$f_a(x_1) = p(x_1)$$

$$f_b(x_2) = p(x_2)$$

## The sum-product algorithm

- The *sum-product* algorithm is an efficient algorithm for exact inference on *tree-structured* graphs
- It is a message passing algorithm as its simpler version for chains
- We will present it on factor graphs, assuming a tree-structured graph giving rise to a factor graph which is a tree
- The algorithm will be applicable to undirected models (i.e. Markov Networks) as well as directed ones (i.e. Bayesian Networks)


## Computing marginals


- We want to compute the marginal probability of  $X$ :

$$p(X) = \sum_{\mathbf{x} \setminus X} p(\mathbf{X})$$

- Generalizing the message passing scheme seen for chains, this can be computed as the product of messages coming from all neighbouring factors  $f_s$ :

$$p(X) = \prod_{f_s \in \text{ne}(X)} \mu_{f_s \rightarrow X}(X)$$






## Factor messages

- Each factor message is the product of messages coming from nodes other than  $X$ , times the factor, summed over all possible values of the factor variables other than  $X$  ( $X_1, \dots, X_M$ ):

$$\mu_{f_s \rightarrow X}(X) = \sum_{X_1} \cdots \sum_{X_M} f_s(X, X_1, \dots, X_M) \prod_{X_m \in \text{ne}(f_s) \setminus X} \mu_{X_m \rightarrow f_s}(X_m)$$






## Node messages

- Each message from node  $X_m$  to factor  $f_s$  is the product of the factor messages to  $X_m$  coming from factors other than  $f_s$ :


$$\mu_{X_m \rightarrow f_s}(X_m) = \prod_{f_l \in \text{ne}(X_m) \setminus f_s} \mu_{f_l \rightarrow X_m}(X_m)$$

## Initialization

- Message passing start from leaves, either factors or nodes
- Messages from leaf factors are initialized to the factor itself (there will be no  $X_m$  different from the destination on which to sum over)



- Messages from leaf nodes are initialized to 1




## Message passing scheme

- The node  $X$  whose marginal has to be computed is designed as root.
- Messages are sent from all leaves to their neighbours
- Each internal node sends its message towards the root as soon as it received messages from all other neighbours
- Once the root has collected all messages, the marginal can be computed as the product of them

## Full message passing scheme

- In order to be able to compute marginals for any node, messages need to pass in all directions:
  - 1 Choose an arbitrary node as root
  - 2 Collect messages for the root starting from leaves
  - 3 Send messages from the root down to the leaves
- All messages passed in all directions using only twice the number of computations used for a single marginal


# Inference example



Consider the joint distribution as product of factors


$$p(\mathbf{X}) = f_a(X_1, X_2)f_b(X_2, X_3)f_c(X_2, X_4)$$

# Inference example



Choose  $X_3$  as root

# Inference example




Send initial messages from leaves

$$\mu_{x_1 \rightarrow f_a}(x_1) = 1$$

$$\mu_{x_4 \rightarrow f_c}(x_4) = 1$$

# Inference example




Send messages from factor nodes to  $X_2$

$$\mu_{f_a \rightarrow X_2}(X_2) = \sum_{X_1} f_a(X_1, X_2)$$

$$\mu_{f_c \rightarrow X_2}(X_2) = \sum_{X_4} f_c(X_2, X_4)$$




# Inference example



Send message from  $X_2$  to factor node  $f_b$

$$\mu_{X_2 \rightarrow f_b}(X_2) = \mu_{f_a \rightarrow X_2}(X_2) \mu_{f_c \rightarrow X_2}(X_2)$$


# Inference example



Send message from  $f_b$  to  $x_3$

$$\mu_{f_b \rightarrow x_3}(x_3) = \sum_{x_2} f_b(x_2, x_3) \mu_{x_2 \rightarrow f_b}(x_2)$$


# Inference example



Send message from root  $x_3$

$$\mu_{x_3 \rightarrow f_b}(x_3) = 1$$


# Inference example



Send message from  $f_b$  to  $x_2$

$$\mu_{f_b \rightarrow x_2}(x_2) = \sum_{x_3} f_b(x_2, x_3)$$

# Inference example




Send messages from  $X_2$  to factor nodes

$$\mu_{X_2 \rightarrow f_a}(X_2) = \mu_{f_b \rightarrow X_2}(X_2) \mu_{f_c \rightarrow X_2}(X_2)$$

$$\mu_{X_2 \rightarrow f_c}(X_2) = \mu_{f_b \rightarrow X_2}(X_2) \mu_{f_a \rightarrow X_2}(X_2)$$

# Inference example




Send messages from factor nodes to leaves

$$\mu_{f_a \rightarrow x_1}(x_1) = \sum_{x_2} f_a(x_1, x_2) \mu_{x_2 \rightarrow f_a}(x_2)$$

$$\mu_{f_c \rightarrow x_4}(x_4) = \sum_{x_2} f_c(x_2, x_4) \mu_{x_2 \rightarrow f_c}(x_2)$$

# Inference example



Compute for instance the marginal for  $X_2$



$$\begin{aligned} p(X_2) &= \mu_{f_a \rightarrow X_2}(X_2) \mu_{f_b \rightarrow X_2}(X_2) \mu_{f_c \rightarrow X_2}(X_2) \\ &= \left[ \sum_{X_1} f_a(X_1, X_2) \right] \left[ \sum_{X_3} f_b(X_2, X_3) \right] \left[ \sum_{X_4} f_c(X_2, X_4) \right] \\ &= \sum_{X_1} \sum_{X_3} \sum_{X_4} f_a(X_1, X_2) f_b(X_2, X_3) f_c(X_2, X_4) \\ &= \sum_{X_1} \sum_{X_3} \sum_{X_4} p(\mathbf{X}) \end{aligned}$$

## Adding evidence

- If some nodes  $\mathbf{X}_e$  are observed, we simply use their observed values instead of summing over all possible values when computing their messages
- After normalization, this gives the conditional probability given the evidence





# Inference example



## Bayesian network

- Take a Bayesian network
- Build a factor graph representing it
- Compute the marginal for a variable (e.g.  $B$ )

# Inference example





Compute the marginal for  $B$

- Leaf factor nodes send messages:

$$\mu_{f_A \rightarrow A} = P(A)$$

$$\mu_{f_D \rightarrow D} = P(D)$$

# Inference example





Compute the marginal for  $B$

- $A$  and  $D$  send messages:

$$\mu_{A \rightarrow f_{A,B,C}}(A) = \mu_{f_A \rightarrow A} = P(A)$$

$$\mu_{D \rightarrow f_{C,D}}(D) = \mu_{f_D \rightarrow D} = P(D)$$

# Inference example





Compute the marginal for  $B$

- $f_{C,D}$  sends message:

$$\mu_{f_{C,D} \rightarrow C}(C) = \sum_D P(C|D) \mu_{f_D \rightarrow D} = \sum_D P(C|D) P(D)$$

# Inference example





Compute the marginal for  $B$

- $C$  sends message:

$$\mu_{C \rightarrow f_{A,B,C}}(C) = \mu_{f_{C,D} \rightarrow C}(C) = \sum_D P(C|D)P(D)$$

# Inference example





## Compute the marginal for $B$

- $f_{A,B,C}$  sends message:

$$\begin{aligned}\mu_{f_{A,B,C} \rightarrow B}(B) &= \sum_A \sum_C P(B|A, C) \mu_{C \rightarrow f_{A,B,C}}(C) \mu_{A \rightarrow f_{A,B,C}}(A) \\ &= \sum_A \sum_C P(B|A, C) P(A) \sum_D P(C|D) P(D)\end{aligned}$$

# Inference example



## Compute the marginal for $B$

- The desired marginal is obtained:

$$\begin{aligned} P(B) &= \mu_{f_{A,B,C \rightarrow B}}(B) = \sum_A \sum_C P(B|A, C)P(A) \sum_D P(C|D)P(D) \\ &= \sum_A \sum_C \sum_D P(B|A, C)P(A)P(C|D)P(D) \\ &= \sum_A \sum_C \sum_D P(A, B, C, D) \end{aligned}$$

## Finding the most probable configuration

- Given a joint probability distribution  $p(\mathbf{X})$
- We wish to find the configuration for variables  $\mathbf{X}$  having the highest probability:

$$\mathbf{X}^{\max} = \underset{\mathbf{X}}{\operatorname{argmax}} p(\mathbf{X})$$

for which the probability is:

$$p(\mathbf{X}^{\max}) = \max_{\mathbf{X}} p(\mathbf{X})$$

## Note

- We want the configuration which is *jointly* maximal for all variables
- We cannot simply compute  $p(X_i)$  for each  $i$  (using the sum-product algorithm) and maximize it



## The max-product algorithm

$$p(\mathbf{X}^{\max}) = \max_{\mathbf{X}} p(\mathbf{X}) = \max_{X_1} \cdots \max_{X_M} p(\mathbf{X})$$

- As for the sum-product algorithm, we can exploit the distribution factorization to efficiently compute the maximum
- It suffices to replace sum with max in the sum-product algorithm

## Linear chain

$$\begin{aligned} \max_{\mathbf{X}} p(\mathbf{X}) &= \max_{X_1} \cdots \max_{X_N} [p(X_1)p(X_2|X_1) \cdots p(X_N|X_{N-1})] \\ &= \max_{X_1} \left[ p(X_1)p(X_2|X_1) \left[ \cdots \max_{X_N} p(X_N|X_{N-1}) \right] \right] \end{aligned}$$

## Message passing

- As for the sum-product algorithm, the max-product can be seen as message passing over the graph.
- The algorithm is thus easily applied to tree-structured graphs via their factor trees:

$$\mu_{f \rightarrow X}(X) = \max_{X_1, \dots, X_M} \left[ f(X, X_1, \dots, X_M) \prod_{X_m \in ne(f) \setminus X} \mu_{X_m \rightarrow f}(X_m) \right]$$
$$\mu_{X \rightarrow f}(X) = \prod_{f_l \in ne(X) \setminus f} \mu_{f_l \rightarrow X}(X)$$

## Recovering maximal configuration

- Messages are passed from leaves to an arbitrarily chosen root  $X_r$
- The probability of maximal configuration is readily obtained as:

$$p(\mathbf{X}^{\max}) = \max_{X_r} \left[ \prod_{f_l \in ne(X_r)} \mu_{f_l \rightarrow X_r}(X_r) \right]$$

- The maximal configuration for the root is obtained as:

$$X_r^{\max} = \operatorname{argmax}_{X_r} \left[ \prod_{f_l \in ne(X_r)} \mu_{f_l \rightarrow X_r}(X_r) \right]$$

- We need to recover maximal configuration for the other variables

## Recovering maximal configuration

- When sending a message towards  $x$ , each factor node should store the configuration of the other variables which gave the maximum:

$$\phi_{f \rightarrow X}(X) = \operatorname{argmax}_{X_1, \dots, X_M} \left[ f(X, X_1, \dots, X_M) \prod_{X_m \in \text{ne}(f) \setminus X} \mu_{X_m \rightarrow f}(X_m) \right]$$

- When the maximal configuration for the root node  $X_r$  has been obtained, it can be used to retrieve the maximal configuration for the variables in neighbouring factors from:

$$X_1^{\max}, \dots, X_M^{\max} = \phi_{f \rightarrow X_r}(X_r^{\max})$$

- The procedure can be repeated *back-tracking* to the leaves, retrieving maximal values for all variables

# Recovering maximal configuration

## Example for linear chain

$$X_N^{\max} = \operatorname{argmax}_{X_N} \mu_{f_{N-1,N} \rightarrow X_N}(X_N)$$


$$X_{N-1}^{\max} = \phi_{f_{N-1,N} \rightarrow X_N}(X_N^{\max})$$

$$X_{N-2}^{\max} = \phi_{f_{N-2,N-1} \rightarrow X_{N-1}}(X_{N-1}^{\max})$$

$$\vdots$$

$$X_1^{\max} = \phi_{f_{1,2} \rightarrow X_2}(X_2^{\max})$$

# Recovering maximal configuration



## Trellis for linear chain

- A *trellis* or *lattice* diagram shows the  $K$  possible states of each variable  $X_n$  one per row
- For each state  $k$  of a variable  $X_n$ ,  $\phi_{f_{n-1,n} \rightarrow X_n}(X_n)$  defines a unique (maximal) previous state, linked by an edge in the diagram
- Once the maximal state for the last variable  $X_N$  is chosen, the maximal states for other variables are recovering following the edges backward.

## Underflow issues

- The max-product algorithm relies on products (no summation)
- Products of many small probabilities can lead to underflow problems
- This can be addressed computing the logarithm of the probability instead
- The logarithm is monotonic, thus the proper maximal configuration is recovered:

$$\log \left( \max_{\mathbf{X}} p(\mathbf{X}) \right) = \max_{\mathbf{X}} \log p(\mathbf{X})$$

- The effect is replacing products with sums (of logs) in the max-product algorithm, giving the *max-sum* one

## Exact inference on general graphs

- The sum-product and max-product algorithms can be applied to tree-structured graphs
- Many applications require graphs with (undirected) loops
- An extension of this algorithms to generic graphs can be achieved with the *junction tree algorithm*
- The algorithm does not work on factor graphs, but on *junction trees*, tree-structured graphs with nodes containing clusters of variables of the original graph
- A message passing scheme analogous to the sum-product and max-product algorithms is run on the junction tree

## Problem

- The complexity on the algorithm is exponential on the maximal number of variables in a cluster, making it intractable for large complex graphs.



## Approximate inference

- In cases in which exact inference is intractable, we resort to *approximate* inference techniques
- A number of techniques for approximate inference exist:
  - loopy belief propagation** message passing on the original graph even if it contains loops
  - variational methods** deterministic approximations, assuming the posterior probability (given the evidence) factorizes in a particular way
  - sampling methods** approximate posterior is obtained sampling from the network

## Loopy belief propagation

- Apply sum-product algorithm even if it is not guaranteed to provide an exact solution
- We assume all nodes are in condition of sending messages (i.e. they already received a constant 1 message from all neighbours)
- A *message passing schedule* is chosen in order to decide which nodes start sending messages (e.g. *flooding*, all nodes send messages in all directions at each time step)
- Information flows many times around the graph (because of the loops), each message on a link replaces the previous one and is only based on the most recent messages received from the other neighbours
- The algorithm can eventually converge (no more changes in messages passing through any link) depending on the specific model over which it is applied