

Chap. 17

Making Complex Decisions

*Decision making methods of what to do today,
given that we may decide again tomorrow
– sequential decision problem.*

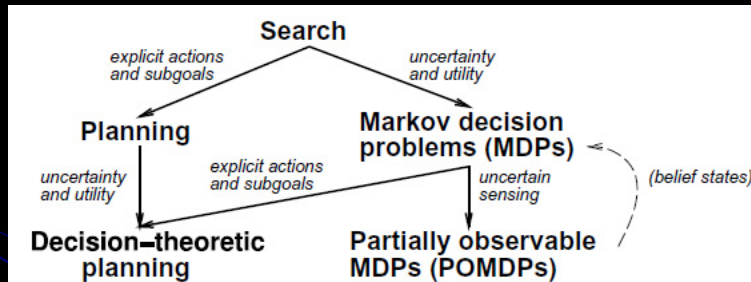
1

Outline

- Sequential Decision problems
 - The agent's utility depends on a sequence of decisions.
 - \supset utilities, uncertainty, sensing.
 - A stochastic generalization of search/planning problem.
 - How sequential decision problem is defined?
- Value iteration
 - How to solve them to produce *optimal behavior* that balances the risks and rewards of acting in an uncertain environment
 - i.e. how to find an optimal policy.
- Policy iteration
 - An alternative way to find optimal policies.

2


Sequential decision problems



	Deterministic	Stochastic
Fully Observable	A*, DFS, BFS, etc.	MDP
Partially Observable		POMDP

3

Markov Decision Process(MDP)




- The specification of a *sequential decision problem* for a *fully observable environment* with a *Markovian Transition Model* and additive *Rewards*.
- Components: a set of states S (initial state= s_0), a set of ACTIONS(s), $\forall s \in S$, *transition model* M_{ij}^a , *reward function* $R(i)$
 - Model $M_{ij}^a \equiv P(j|i, a)$ = probability that doing a in i leads to j : Markov transition
 - Utility function will depend on a *sequence of states* (i.e. an environment history) rather than a single state.
 - Each state has a *reward* $R(i)$
 - = -0.04 (small penalty) for nonterminal states; = ± 1 for terminal states
 - Utility of an environment history \approx the *sum* of the rewards received.

4


Solving MDPs

- In search problems, aim is to find an optimal *sequence of actions*.
- In MDPs, aim is to find an optimal *policy* π
i.e. best action for *every possible state*
(because can't predict where one will end up)
 $\pi(s)$: the action recommended by the policy π for state s .
- The quality (i.e. value) of a policy is measured by the *expected utility* of the possible environment histories generated by that policy.
- An *optimal policy* π^* : a policy with the *highest expected utility (MEU)*.
- Given π^* , the agent decides what to do by consulting its current percept, which tells it the current state s , and then executing the action $\pi^*(s)$.
- The balance of risk and reward changes depends on the value of $R(s)$ for the nonterminal states.
- The careful *balancing of risk & reward* is a characteristic of MDPs which doesn't arise in deterministic search problem but in real-world decision problem.

5



An optimal policy for the stochastic environment with $R(s) = -0.04$ in the nonterminal states



Optimal policies for four different ranges of $R(s)$.

6

Utility

- The performance of agent is measured by a sum or rewards for the states visited.
- The choices for the utility function on environment histories.
- Is there a finite horizon or an infinite horizon for decision making?
 - $U_h([s_0, s_1, \dots, s_{N+k}]) = U_h([s_0, s_1, \dots, s_N])$, for all $k > 0$.
 - With a finite horizon (with a fixed time N), the optimal action in a given state could change over time.
 - the optimal policy for a finite horizon is non-stationary.
 - With an infinite horizon, the optimal action depends only on the current state and is stationary.
 - Policies for the infinite horizon (i.e. no fixed deadline) case are simpler.
- How to calculate the utility of state sequences?
 - A question in multiattribute utility theory where s_i is viewed as attribute of the state sequence $[s_0, s_1, \dots]$.
 - Assumption: the agent's preferences b/t state sequences are stationary:

7

Utility: cont.

- In *sequential* decision problems, preferences are expressed between *sequences* of states.
- *Preference-independence* Assumption: agent's preferences between state sequences are *stationary*.
 - If $s_0 = s'_0$ in two sequence $[s_0, s_1, s_2, \dots]$ and $[s'_0, s'_1, s'_2, \dots]$, then two sequences should be preference-ordered the same way as $[s_1, s_2, \dots]$ and $[s'_1, s'_2, \dots]$.
- *Additive Rewards*: Usually uses an *additive* utility function

$$U_h([s_0, s_1, \dots, s_n]) = R(s_0) + R(s_1) + R(s_2) + \dots + R(s_n)$$

(cf. path cost in search problems).
- *Discounted Rewards*:

$$U_h([s_0, s_1, \dots, s_n]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots + \gamma^n R(s_n)$$

$$= \sum_{t=0}^n \gamma^t R(s_t).$$

where γ : discount factor $\in [0, 1]$,
 the preference of an agent for current rewards over future rewards.
 $\gamma \rightarrow 0$, rewards in the distant future are insignificant; $\gamma = 1$, additive rewards.

Utility: cont.

- If the environment contains no terminal state or if the agent never reaches one
→ infinite horizon.
- Solutions for some problems with infinite horizons:
 - Utility of an infinite sequence is finite with discounted rewards:

$$U_\gamma([s_0, s_1, \dots, s_n]) = \sum_{t=0}^{\infty} \gamma^t R(s_t) \leq \sum_{t=0}^{\infty} \gamma^t R_{max} = R_{max} / (1 - \gamma).$$
 - If the environment contains the terminal states and if the agent is guaranteed to get to one eventually, then no need to compare infinite sequences.
 - A *Proper policy*: a policy that is guaranteed to reach a terminal state.
 - An improper policy can cause the standard algorithms for solving MDPs to fail with additive rewards – good to use discounted rewards.
 - Compare infinite sequences in terms of the average reward per time step.

9

Utilities : cont.

- Compare policies, (π_1, π_2) , by their EU when executing them: U^{π_1}, U^{π_2} .
- *Utility of a state* (a.k.a. its value) : the EU by executing a policy π in s

$$U^\pi(s) = E [\sum_{t=0}^{\infty} \gamma^t R(s_t) \mid \pi, s_0=s]$$

where the expectation is w.r.t. the probability distribution over state sequence decided by s and π .
- How to choose b/t policies?
 - A given policy generates a whole range of possible state sequences
 - The *value of a policy* is the *Expected Sum* of discounted rewards obtained where the expectation is taken over all possible state sequences that could occur: *Expected Utility of a policy*.
 - An *optimal policy* : $\pi_s^* = \operatorname{argmax}_\pi U^\pi(s) = \operatorname{argmax}_\pi E[\sum_{t=0}^{\infty} \gamma^t R(s_t) \mid \pi]$
where s is the start state, s_t is the state the agent is in after executing π for t steps.
 - The optimal policy is independent of the starting state, using the discounted utilities with infinite horizons.

10

Utility: cont.

- True Utility of a state (a.k.a. its value):
 $U_h(s) = \text{expected sum of discounted rewards until termination under the optimal policy} = U^{\pi^*}(s)$
- cf) $R(s)$: the short-term reward for being in s
 $U(s)$: the long-term total reward from s onwards (in a sequence).
- Given the utilities of the states $U(s)$, choosing the best action is just MEU: choose the action such that the EU of the immediate successors is highest:
 $\pi^*(i) = \operatorname{argmax}_a \sum_j U(j) M_{ij}^a$ where $M_{ij}^a = P(j|i, a)$, $a \in A(i)$
- Optimal policy and state values ($U(s)$) for the given $R(s) = -.04$:

3	→	→	→	+1	3	0.812	0.868	0.912	+1
2	↑		↑	-1	2	0.762		0.660	-1
1	↑	←	←	←	1	0.705	0.655	0.611	0.388
	1	2	3	4		1	2	3	4

11

Bellman equation

- How to find the optimal policy?
- **Value Iteration Algorithm**: Calculate the utility of each state and use the state utilities to select an optimal action in each state.
- *the utility of a state* = the immediate reward for that state
 + the expected discounted utility of the next state assuming that the agent chooses the optimal action.
- **Bellman equation (1957)**:

$$U(i) = R(i) + \gamma \max_a \sum_j U(j) M_{ij}^a$$

$$U(1,1) = -.04$$

$$+ \max \begin{cases} 0.8U(1,2) + 0.1U(2,1) + 0.1U(1,1), & (\text{Up}) \\ 0.9U(1,1) + 0.1U(1,2), & (\text{Left}) \\ 0.9U(1,1) + 0.1U(2,1), & (\text{Down}) \\ 0.8U(2,1) + 0.1U(1,2) + 0.1U(1,1) \} & (\text{Right}) \end{cases}$$

One equation per state = n (nonlinear) equation in n unknowns

12

Value Iteration Algorithm

- Iterative approach using n Bellman equations for n possible states.
- Idea: Start with arbitrary utility values
 - Update to make them **locally consistent** with Bellman eqn.
 - Everywhere locally consistent \Rightarrow global optimality
 - repeat until “equilibrium”

$$U_k(i) \leftarrow R(i) + \gamma \max_a \sum_j U_k(j) M_{ij}^a \quad \text{for all } i, \text{ at step } k.$$


-- a unique solution to Bellman equation. \Rightarrow optimal policy

-- a propagation of information through the state space b.m.o. local updates.

```

function VALUE-ITERATION( $mdp, \epsilon$ ) returns a utility function
inputs:  $mdp$ , an MDP with states  $S$ , actions  $A(s)$ , transition model  $P(s' | s, a)$ ,
       rewards  $R(s)$ , discount  $\gamma$ 
 $\epsilon$ , the maximum error allowed in the utility of any state
local variables:  $U, U'$ , vectors of utilities for states in  $S$ , initially zero
                  $\delta$ , the maximum change in the utility of any state in an iteration

repeat
     $U \leftarrow U'$ ;  $\delta \leftarrow 0$ 
    for each state  $s$  in  $S$  do
         $U'[s] \leftarrow R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$ 
        if  $|U'[s] - U[s]| > \delta$  then  $\delta \leftarrow |U'[s] - U[s]|$ 
    until  $\delta < \epsilon(1 - \gamma)/\gamma$ 
return  $U$ 
    
```



3

Policy iteration (Howard, 1960)

- Idea: search for optimal policy and utility values simultaneously
- **Policy Evaluation**
 - Given a policy π_k , calculate $U_k = U^{\pi_k}$, the utility of each state if π_k were to be executed.

+ Policy Improvement

Calculate a new MEU policy π_{k+1} , using one-step look-ahead based on U_k .

- Algorithm:

$\pi_1 \leftarrow$ an arbitrary initial policy
 repeat until no change in π
 compute utilities U_i given π_i
 update π_i as if utilities were correct (i.e. local MEU)

- To compute utilities given a fixed π

$$U(i) = R(i) + \gamma \sum_j U(j) M_{ij}^{\pi(i)} \quad \text{for all } i \text{ where } M_{ij}^{\pi(i)} = P(j|i, \pi(i)),$$

where $\pi(i)$ specifies the action $\pi(i)$ in state i under the policy π

$$U_k(1,1) = -0.04 + 0.8U_k(1,2) + 0.1U_k(2,1) + 0.1U_k(1,1),$$

$$U_k(1,2) = -0.04 + 0.8U_k(1,3) + 0.2U_k(1,2),$$

.....

14

Policy iteration (Howard, 1960)

```
function POLICY-ITERATION(mdp) returns a policy
  inputs: mdp, an MDP with states S, actions A(s), transition model  $P(s' | s, a)$ 
  local variables: U, a vector of utilities for states in S, initially zero
                   $\pi$ , a policy vector indexed by state, initially random

  repeat
     $U \leftarrow \text{POLICY-EVALUATION}(\pi, U, \text{mdp})$ 
    unchanged?  $\leftarrow$  true
    for each state s in S do
      if  $\max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s'] > \sum_{s'} P(s' | s, \pi[s]) U[s']$  then do
         $\pi[s] \leftarrow \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$ 
    unchanged?  $\leftarrow$  false
  until unchanged?
  return  $\pi$ 
```

- n simultaneous *linear* equations in n unknowns, solve in $O(n^3)$ by standard linear algebra methods.
- The simplified Bellman update: often more efficient

$$U_{k+1}(i) = R(i) + \gamma \sum_j U_k(j) M^{\pi(i)}_{ij} \text{ for all } i \text{ where } M^{\pi(i)}_{ij} = P(j|i, \pi(i)),$$

15